



МИНОБРНАУКИ РОССИИ
Федеральное государственное бюджетное образовательное учреждение высшего
образования
«МИРЭА – Российский технологический университет»
РТУ МИРЭА

Институт кибербезопасности и цифровых технологий
КБ-4 «Интеллектуальные системы информационной безопасности»

Отчет по практической и лабораторной работе
по дисциплине: «Анализ защищенности систем искусственного
интеллекта»

Выполнил:
Тимофеев И.О
ББМО-01-22

Проверил:
К.т.н. Спирин А.А.

Москва, 2023

Выполним импорт необходимых библиотек.

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms, datasets
```

Загрузим набор данных (MNIST), разобьем данные на подвыборки.

```
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.0, ), (1.0, ))])
dataset = datasets.MNIST(root = './data', train=True, transform = transform, download=True)
train_set, val_set = torch.utils.data.random_split(dataset, [50000, 10000])
test_set = datasets.MNIST(root = './data', train=False, transform = transform, download=True)
train_loader = torch.utils.data.DataLoader(train_set, batch_size=1, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=1, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=1, shuffle=True)
print("Training data:", len(train_loader), "Validation data:", len(val_loader), "Test data:", len(test_loader))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|#####| 9912422/9912422 [00:00<00:00, 118647742.08it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|#####| 28881/28881 [00:00<00:00, 47993539.55it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw
```

Настроим использование графического ускорителя.

```
[3] use_cuda=True
    device = torch.device("cuda" if (use_cuda and torch.cuda.is_available()) else "cpu")
```

1 Создание атак на модель НС

Создадим класс НС на основе фреймворка torch

```
[4] class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)
    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output
```

Проверим работоспособность созданного класса НС.

```
[5] model = Net().to(device)
```

Создадим оптимизатор, функцию потерь и трейнер сети.

```
[6] optimizer = optim.Adam(model.parameters(), lr=0.0001, betas=(0.9, 0.999))
    criterion = nn.NLLLoss()
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.1, patience=3)
```

Определим функцию обучения сети.

```
[7] def fit(model,device,train_loader,val_loader,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss
```


Обучим модель.

```
[8] loss, val_loss = fit(model, device, train_loader, val_loader, 10)

Fitting the model...
/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py:1345: UserWarning:
  warnings.warn(warn_msg)
Epoch: 1 Loss: 0.27364798618042624 Val_Loss: 0.1514760968370528
Epoch: 2 Loss: 0.1085534713517188 Val_Loss: 0.10566764865441187
Epoch: 3 Loss: 0.08270403996756964 Val_Loss: 0.09490188534468656
Epoch: 4 Loss: 0.07184107978025048 Val_Loss: 0.08452242758264042
Epoch: 5 Loss: 0.06318974395321775 Val_Loss: 0.08870745507981724
Epoch: 6 Loss: 0.06183685142550104 Val_Loss: 0.08463778928197853
Epoch: 7 Loss: 0.05919382139317826 Val_Loss: 0.09058659238292247
Epoch: 8 Loss: 0.05746152144902565 Val_Loss: 0.0884336885243108
Epoch: 9 Loss: 0.04317616631293602 Val_Loss: 0.07717185551352328
Epoch: 10 Loss: 0.04143624664735127 Val_Loss: 0.08422930549846026
```

Построим графики потерь при обучении и валидации в зависимости от эпохи.

```
[9] fig = plt.figure(figsize=(5,5))
plt.plot(np.arange(1,11), loss, "*-",label="Loss")
plt.plot(np.arange(1,11), val_loss,"o-",label="Val Loss")
plt.xlabel("Num of epochs")
plt.legend()
plt.show()
```



Создадим функции атак FGSM, I-FGSM, MI-FGSM.

```
[10] def fgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

[11] def ifgsm_attack(input,epsilon,data_grad):
    iter = 10
    alpha = epsilon/iter
    pert_out = input
    for i in range(iter-1):
        pert_out = pert_out + alpha*data_grad.sign()
        pert_out = torch.clamp(pert_out, 0, 1)
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:
            break
    return pert_out
```

```
[12] def mifgsm_attack(input,epsilon,data_grad):
    iter=10
    decay_factor=1.0
    pert_out = input
    alpha = epsilon/iter
    g=0
    for i in range(iter-1):
        g = decay_factor*g + data_grad/torch.norm(data_grad,p=1)
        pert_out = pert_out + alpha*torch.sign(g)
        pert_out = torch.clamp(pert_out, 0, 1)
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:
            break
    return pert_out
```


Создадим функцию проверки.

```
[13] def test(model,device,test_loader,epsilon,attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data,epsilon,data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data,epsilon,data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
        if (epsilon == 0) and (len(adv_examples) < 5):
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))
    return final_acc, adv_examples
```


Построим графики успешности атак(Ассигуасы/эпсилон) и примеры выполненных атак в зависимости от степени возмущения epsilon.

```
epsilons = [0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
for attack in ("fgsm","ifgsm","mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model, device, test_loader, eps, attack)
        accuracies.append(acc)
        examples.append(ex)
    plt.figure(figsize=(5,5))
    plt.plot(epsilons, accuracies, "*-")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()
    cnt = 0
    plt.figure(figsize=(8,10))
    for i in range(len(epsilons)):
        for j in range(len(examples[i])):
            cnt += 1
            plt.subplot(len(epsilons), len(examples[0]), cnt)
            plt.xticks([], [])
            plt.yticks([], [])
            if j == 0:
                plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
            orig, adv, ex = examples[i][j]
            plt.title("{} -> {}".format(orig, adv))
            plt.imshow(ex, cmap="gray")
    plt.tight_layout()
    plt.show()
```

Epsilon: 0	Test Accuracy = 9761 / 10000 = 0.9761
Epsilon: 0.007	Test Accuracy = 9721 / 10000 = 0.9721
Epsilon: 0.01	Test Accuracy = 9730 / 10000 = 0.973
Epsilon: 0.02	Test Accuracy = 9665 / 10000 = 0.9665
Epsilon: 0.03	Test Accuracy = 9624 / 10000 = 0.9624
Epsilon: 0.05	Test Accuracy = 9423 / 10000 = 0.9423
Epsilon: 0.1	Test Accuracy = 8680 / 10000 = 0.868
Epsilon: 0.2	Test Accuracy = 5786 / 10000 = 0.5786
Epsilon: 0.3	Test Accuracy = 2929 / 10000 = 0.2929




	1 -> 1	3 -> 3	4 -> 4	8 -> 8	3 -> 3
Eps: 0					
Eps: 0.007	0 -> 0 	1 -> 1 	8 -> 8 	7 -> 7 	6 -> 6
Eps: 0.01	6 -> 0 	2 -> 2 	2 -> 2 	0 -> 0 	6 -> 6
Eps: 0.02	4 -> 4 	5 -> 5 	1 -> 1 	7 -> 7 	2 -> 2
Eps: 0.03	4 -> 4 	8 -> 8 	2 -> 2 	9 -> 9 	7 -> 7
Eps: 0.05	6 -> 6 	5 -> 5 	7 -> 7 	2 -> 2 	3 -> 3
Eps: 0.1	2 -> 2 	5 -> 5 	1 -> 1 	7 -> 7 	0 -> 0




```

Epsilon: 0      Test Accuracy = 9780 / 10000 = 0.978
Epsilon: 0.007  Test Accuracy = 9722 / 10000 = 0.9722
Epsilon: 0.01   Test Accuracy = 9725 / 10000 = 0.9725
Epsilon: 0.02   Test Accuracy = 9687 / 10000 = 0.9687
Epsilon: 0.03   Test Accuracy = 9660 / 10000 = 0.966
Epsilon: 0.05   Test Accuracy = 9510 / 10000 = 0.951
Epsilon: 0.1    Test Accuracy = 8931 / 10000 = 0.8931
Epsilon: 0.2    Test Accuracy = 6417 / 10000 = 0.6417
Epsilon: 0.3    Test Accuracy = 3627 / 10000 = 0.3627

```





2 Защита от атак

Создадим 2 класса НС.

```
class NetF(nn.Module):
    def __init__(self):
        super(NetF, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x
```

```
class NetF1(nn.Module):
    def __init__(self):
        super(NetF1, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3, 1)
        self.conv2 = nn.Conv2d(16, 32, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(4608, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x
```

Переопределим функцию обучения и тестирования.

```
def fit(model,device,optimizer,scheduler,criterion,train_loader,val_loader,Temp,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                output = F.log_softmax(output/Temp,dim=1)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss

def test(model,device,test_loader,epsilon,Temp,attack):
    correct=0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        output = F.log_softmax(output/Temp,dim=1)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data,epsilon,data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data,epsilon,data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
        if (epsilon == 0) and (len(adv_examples) < 5):
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
```

Создадим функцию защиты методом дистилляции.


```
def defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons):
    modelF = NetF().to(device)
    optimizerF = optim.Adam(modelF.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF = optim.lr_scheduler.ReduceLROnPlateau(optimizerF, mode='min', factor=0.1, patience=3)
    modelF1 = NetF1().to(device)
    optimizerF1 = optim.Adam(modelF1.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF1 = optim.lr_scheduler.ReduceLROnPlateau(optimizerF1, mode='min', factor=0.1, patience=3)
    criterion = nn.NLLLoss()
    lossF,val_lossF=fit(modelF,device,optimizerF,schedulerF,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF, "-.",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF,"o-",label="Val Loss")
    plt.title("Network F")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    #converting target labels to soft labels
    for data in train_loader:
        input, label = data[0].to(device),data[1].to(device)
        softlabel = F.log_softmax(modelF(input),dim=1)
        data[1] = softlabel
    lossF1,val_lossF1=fit(modelF1,device,optimizerF1,schedulerF1,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF1, "-.",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF1,"o-",label="Val Loss")
    plt.title("Network F'")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    model = NetF1().to(device)
    model.load_state_dict(modelF1.state_dict())
    for attack in ("fgsm","ifgsm","mifgsm"):
        accuracies = []
        examples = []
        for eps in epsilons:
            acc, ex = test(model,device,test_loader,eps,attack)
            accuracies.append(acc)
            examples.append(ex)
        plt.figure(figsize=(5,5))
        plt.plot(epsilons, accuracies, "-.")
        plt.title(attack)
        plt.xlabel("Epsilon")
        plt.ylabel("Accuracy")
        plt.show()
        cnt = 0
        plt.figure(figsize=(8,10))
        for i in range(len(epsilons)):
            for j in range(len(examples[i])):
                cnt += 1
                plt.subplot(len(epsilons),len(examples[0]),cnt)
                plt.xticks([], [])
                plt.yticks([], [])
                if j == 0:
                    plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
                orig,adv,ex = examples[i][j]
                plt.title("{} -> {}".format(orig, adv))
                plt.imshow(ex, cmap="gray")
        plt.tight_layout()
```

Получим результаты оценки защищенных

```
Temp=100
epochs=10
epsilons=[0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons)
```


Fitting the model...

```
Epoch: 1 Loss: 0.5586636638233529 Val_Loss: 0.00013734494727104903
Epoch: 2 Loss: 0.32000435114445375 Val_Loss: 3.885443676263094e-05
Epoch: 3 Loss: 0.22065452498864171 Val_Loss: 2.6573986560106278e-05
Epoch: 4 Loss: 0.166122699417228 Val_Loss: 2.089180740586016e-06
Epoch: 5 Loss: 0.13697167539701155 Val_Loss: 2.0265563023258437e-10
Epoch: 6 Loss: 0.12113108773361844 Val_Loss: 1.0716508119367063e-05
Epoch: 7 Loss: 0.11046593322182684 Val_Loss: 1.7166009683933225e-09
Epoch: 8 Loss: 0.10070784397399411 Val_Loss: 1.8325391924008726e-06
Epoch: 9 Loss: 0.09728810295630186 Val_Loss: 6.29195295914542e-06
Epoch: 10 Loss: 0.08716456005690097 Val_Loss: 1.969227523659356e-08
```



Fitting the model...



```
Epoch: 1 Loss: 0.6708522244577682 Val_Loss: 8.905845731496811e-05
Epoch: 2 Loss: 0.4460383526137011 Val_Loss: 8.857200965285301e-05
Epoch: 3 Loss: 0.38187247984231215 Val_Loss: 2.3023875430226327e-06
Epoch: 4 Loss: 0.33720084358059105 Val_Loss: 2.4729261049287745e-06
Epoch: 5 Loss: 0.29825992335180374 Val_Loss: 1.5593550726771354e-05
Epoch: 6 Loss: 0.2625956706629719 Val_Loss: 2.759788396360818e-07
Epoch: 7 Loss: 0.23289492150114827 Val_Loss: 3.2116492511704563e-07
Epoch: 8 Loss: 0.20901024173459745 Val_Loss: 1.4200437907129527e-05
Epoch: 9 Loss: 0.18101272500992469 Val_Loss: 0.0003574684738783617
Epoch: 10 Loss: 0.17011370223791447 Val_Loss: 7.655074005015194e-08
```



```

Epsilon: 0      Test Accuracy = 9282 / 10000 = 0.9282
Epsilon: 0.007  Test Accuracy = 9266 / 10000 = 0.9266
Epsilon: 0.01   Test Accuracy = 9253 / 10000 = 0.9253
Epsilon: 0.02   Test Accuracy = 9186 / 10000 = 0.9186
Epsilon: 0.03   Test Accuracy = 9065 / 10000 = 0.9065
Epsilon: 0.05   Test Accuracy = 8850 / 10000 = 0.885
Epsilon: 0.1    Test Accuracy = 7835 / 10000 = 0.7835
Epsilon: 0.2    Test Accuracy = 4814 / 10000 = 0.4814
Epsilon: 0.3    Test Accuracy = 2124 / 10000 = 0.2124
Epsilon: 0      Test Accuracy = 9285 / 10000 = 0.9285
Epsilon: 0.007  Test Accuracy = 9264 / 10000 = 0.9264
Epsilon: 0.01   Test Accuracy = 9239 / 10000 = 0.9239
Epsilon: 0.02   Test Accuracy = 9177 / 10000 = 0.9177
Epsilon: 0.03   Test Accuracy = 9107 / 10000 = 0.9107
Epsilon: 0.05   Test Accuracy = 8819 / 10000 = 0.8819
Epsilon: 0.1    Test Accuracy = 7839 / 10000 = 0.7839
Epsilon: 0.2    Test Accuracy = 4839 / 10000 = 0.4839
Epsilon: 0.3    Test Accuracy = 2130 / 10000 = 0.213
Epsilon: 0      Test Accuracy = 9280 / 10000 = 0.928
Epsilon: 0.007  Test Accuracy = 9262 / 10000 = 0.9262
Epsilon: 0.01   Test Accuracy = 9225 / 10000 = 0.9225
Epsilon: 0.02   Test Accuracy = 9172 / 10000 = 0.9172
Epsilon: 0.03   Test Accuracy = 9056 / 10000 = 0.9056
Epsilon: 0.05   Test Accuracy = 8834 / 10000 = 0.8834
Epsilon: 0.1    Test Accuracy = 7843 / 10000 = 0.7843
Epsilon: 0.2    Test Accuracy = 4825 / 10000 = 0.4825
Epsilon: 0.3    Test Accuracy = 2111 / 10000 = 0.2111

```



Заключение

В данной лабораторной работе по изучению защиты от атак на модели НС, использовался набор данных MNIST. Далее были созданы атаки на модель НС. Созданы функции атак FGSM, I-FGMS, MI-FGSM. Также были созданы графики успешности атак и примеры выполненных атак. Помимо этого, была создан защитный метод дистилляции, а также проведена оценка результата работы. Основная идея защитной дистилляции заключается в обучении устойчивой модели, путем передачи знаний от базовой модели, подверженной атакам, к новой модели, которая спроектирована для устойчивости к различным атакам. Дистилляция дает более плоские локальные минимумы. Следовательно, небольшие изменения во входных данных с меньшей вероятностью изменят прогнозируемые значения.